

Supplements

This material supplements the research article “Mobilising Clinical AI Skills: A Mixed-Methods Evaluation of a National Medical Professional Upskilling Programme”.

S1 Exploratory Cluster Analysis

The exploratory cluster analysis served stage 2 of the type-construction framework (grouping of cases and analysis of empirical regularities); its role was hypothesis generation, not confirmatory classification.

Of the fifteen person-level binary indicators, twelve were used as active variables: medtech industry role, teaching, further AI-related education, structured fellowship programmes, individual research fellowships, the three publication categories (contributions to research, to opinion/debate, and to knowledge), expert speaking, written opinion sharing, policy engagement, and research activities. Three indicators were treated as passive (supplementary) descriptors and did not shape the solution: post-fellowship clinical activity (near-constant at 98%), influencer activity (4%) and med-tech founding (7%). Near-constant and very rare indicators were demoted because the chi-square metric underlying multiple correspondence analysis weights categories inversely to their frequency, so that categories held by only one to four individuals would exert disproportionate leverage on the solution. Cohort membership was likewise treated as passive, because cohorts differ systematically in follow-up time and an active cohort variable would partly cluster observation windows rather than professional profiles.

Multiple correspondence analysis (MCA) was applied to the active indicators. Row coordinates were retained on the leading components explaining at least 80% of total inertia (variant A: 8 of 12 components; variant B: 7 of 11), a deliberately conservative truncation that retains most of the indicator space while removing trailing, noisedominated dimensions. Agglomerative hierarchical clustering with Ward's minimumvariance criterion was then performed on the retained coordinates, and partitions with $k = 2$ to 6 clusters were extracted. The number of clusters was assessed using the average silhouette width.

A sensitivity variant (variant B) merged the expert-speaking, written-opinion-sharing and influencer indicators into a single public-engagement indicator (eleven active variables), testing whether splitting outward-facing communication across three sparse indicators fragmented a coherent trait. As an additional sensitivity analysis, the three comparison dimensions were collapsed into aggregated scores (career pathway 0-5, publications 0-3, activities 0-6), standardised, and clustered with the same procedures, testing whether the binary coding had obscured structure carried by intensity information.

Supplementary Table S1. Codebook: comparison dimensions, categories and person-level binary indicators.

Dimension	Category / indicator	Definition
Career pathway	Clinical (NHS) career	Clinical positions in the NHS, full- or part-time (near-constant; passive descriptor)
Career pathway	Med-tech industry role	Co-founder, associate medical director, clinical AI lead, clinical scientist, medical consultant/advisor roles (part- or full-time)
Career pathway	Teaching	Honorary or substantive lectureships
Career pathway	Further AI-related education	PhD, MSc, MSt or PGDip commenced during or after the fellowship
Career pathway	Structured fellowship programmes	e.g., NHS Topol Digital Health Fellowship, NHS Clinical Entrepreneur Programme (secondary descriptor in the final typology)
Career pathway	Individual research fellowships	e.g., NIHR and other individually held AI-related research fellowships
Publications	Contribution to research	Original research, research protocols, case studies
Publications	Contribution to opinion/debate	Debate articles, editorials, commentaries, perspectives, letters
Publications	Contribution to knowledge	Reviews (incl. narrative and systematic), book chapters
Publications	Excluded	Preprints, conference articles, conference abstracts, posters
AI-related activities	Expert speaking	Speaker at congresses, symposia, webinars, panels, podcasts, YouTube
AI-related activities	Written opinion sharing	Interviews, blog articles
AI-related activities	Policy engagement	Committee membership, delegate and advisory roles
AI-related activities	Influencer activity	Podcast hosting, blog editorship (passive descriptor)
AI-related activities	Research activities	Academic conferences, major research collaborations
AI-related activities	Med-tech founding	Founding an AI company (passive descriptor)

Publications and AI-related activities were captured from each individual's fellowship year until 31 January 2026, career information until 9 February 2026. Each alumnus was assigned one binary indicator per category, denoting the presence of at least one item.


S2 Stability Assessment

Robustness of every candidate partition was assessed by four prespecified procedures. First, cluster-wise bootstrap stability following Hennig's clusterboot scheme: 500 bootstrap resamples of the 57 alumni were drawn with replacement; the entire pipeline (MCA, component truncation, Ward clustering, cut at k) was re-run on each resample; for each original cluster, the maximal Jaccard similarity with any resample cluster was computed over the individuals present in the resample and averaged across resamples. Values of at least 0.75 indicate stable clusters, values of 0.5 or below indicate dissolved clusters.

Second, leave-one-out sensitivity: for each alumnus, the pipeline was re-run on the remaining 56 individuals and agreement with the full solution was quantified by the adjusted Rand index (ARI) over the shared individuals. This diagnostic addresses the specific fragility of small samples, in which single individuals can alter the merge order of hierarchical algorithms.

Third, a cross-method comparison: partitioning around medoids (PAM; 20 random initialisations, iterative medoid update) was run on simple-matching distances computed from the untruncated binary indicators, and its agreement with the MCA-based solution at the same k was quantified by the ARI. Simple matching (in which joint absence counts as agreement) was chosen deliberately: the indicators are outcome markers within a defined graduate cohort, in which shared absence of public activity constitutes genuine resemblance; under a Jaccard coefficient, the largest emerging pattern, defined by joint zeros, could not form by construction. Because this check varies the representation, the metric and the algorithm simultaneously, disagreement with the primary solution is highly informative.

Fourth, a permutation test: 200 null datasets were generated by independently permuting each indicator column, preserving all marginal frequencies while destroying the joint structure; for each null dataset the maximal average silhouette across $k = 2-6$ was computed and compared with the observed value.




Supplementary Figs. S1

Supplementary Table S3. Stability metrics of the exploratory cluster analysis.

Variant	Silhouette by k	Selected k	Bootstrap Jaccard (per cluster)	Leave-one-out ARI	ARI vs PAM	Perm. p
Variant A (12 active)	k=2: 0.185; k=3: 0.212; k=4: 0.186; k=5: 0.212; k=6: 0.224	k=6; sizes 20/5/10/6/3/13	0.63, 0.61, 0.42, 0.71, 0.49, 0.47	mean 0.83, min 0.43	0.37	0.030

Variant B (11 active)	k=2: 0.217; k=3: 0.187; k=4: 0.225; k=5: 0.247; k=6: 0.27	k=6; sizes 6/6/6/25/8/6	0.68, 0.52, 0.62, 0.74, 0.53, 0.48	mean 0.88, min 0.54	0.22	—
-----------------------	---	----------------------------	------------------------------------	---------------------	------	---


Bootstrap Jaccard: cluster-wise stability over 500 resamples of the full pipeline; values ≥ 0.75 indicate stable clusters (Hennig). Leave-one-out ARI: adjusted Rand index between the full solution and each $n-1$ solution. ARI vs PAM: agreement with partitioning around medoids on simple-matching distances. Permutation p : observed maximal silhouette against 200 null datasets with preserved marginal frequencies (computed for variant A; an analogous test on aggregated dimension scores yielded $p = 0.090$). No partition met the stability criteria.



Supplementary Figs. S2

No partition met these criteria (**Supplementary Table S3; Supplementary Figs. S1 and S2**). Cluster-wise bootstrap Jaccard values ranged from 0.42 to 0.74 across both variants and all k ; leave-one-out ARI reached minima of 0.43 (variant A) and 0.54 (variant B); crossmethod agreement was low (ARI 0.37 and 0.22 at $k = 6$); and the observed clustering quality only marginally exceeded the permutation null (variant A: observed maximal silhouette 0.224 against a null 95th percentile of 0.220, $p = 0.030$; aggregated scores: $p = 0.090$). The aggregated-score sensitivity analysis reproduced the same verdict (maximal silhouette 0.36; bootstrap Jaccard predominantly below 0.7). Together with the census design, these results indicate that alumni are distributed along a continuum of engagement intensity (**Supplementary Fig. S3**) rather than in well-separated groups. Consistent with the exploratory role of stage 2, the recurrent attribute combinations of

the best-scoring solution were therefore carried forward solely as candidate groupings for the qualitative type construction; no statistical typology was claimed. All resampling procedures used fixed random seeds; results are exactly reproducible from the analysis code.



Supplementary Figs. S3

S3 Rule-based Classification Procedure

Following qualitative type construction (stage 3), each type was operationalised through an explicit assignment rule over the binary indicators (**Supplementary Table S2**). The rules constitute the reduction of the attribute space in the sense of Lazarsfeld and Barton: they re-describe the constructed types transparently and make the assignment reproducible; they are not an externally validated classifier.

Supplementary Table S2. Assignment rules and rule matches (n = 57).

Type	Assignment rule	Rule matches
Embedded Clinician	Residual rule: no other rule matched AND at most 2 active indicators present	19
Academic Researcher	Research activities = 1; OR contribution to research = 1 AND (further AI education = 1 OR individual research fellowship = 1)	16
Public Communicator	(Expert speaking OR written opinion sharing OR influencer activity = 1) AND (contribution to knowledge OR teaching OR contribution to opinion/debate = 1); OR (contribution to opinion/debate = 1 AND contribution to knowledge = 1)	18
Policy Shaper	Policy engagement = 1	13
Industry Mover	Med-tech industry role = 1	9

Rules were applied without a precedence order. 41 alumni matched exactly one rule and were assigned directly; 16 records matched several rules, or none, and were assigned by coder consensus (Table S4). The Embedded Clinician rule is residual by design and evaluated only where no other rule matched.

Rules were applied to all 57 records without a precedence order. A fixed precedence (for example, industry before policy) would have encoded an arbitrary hierarchy among the types' logics of influence; instead, every record matching more than one rule was referred to adjudication. The Embedded Clinician rule is residual by design: it was evaluated only for records matching no other rule, and additionally required at most two

active indicators. The cap prevents records with a moderately active but rule-atypical profile from being classified as clinically embedded by default; the single such record was referred to adjudication.


Of the 57 records, 41 matched exactly one rule and were assigned directly. The remaining 16 records (15 multi-rule, 1 no-rule) were reviewed independently by both coders (KW, SP) and resolved in consensus meetings. Adjudication followed the dominant logic of influence in the full footprint record (market, institutions, scientific evidence, public discourse, or clinical practice), supported by combination-specific guiding questions (for example, for policy/communicator overlaps: whether the record is dominated by institutional mandates or by public-facing activity). Each consensus decision was documented with a short written rationale, forming an audit trail (**Supplementary Table S4** summarises combinations and outcomes).

Supplementary Table S4. Consensus adjudication of the 16 multi-rule and no-rule records.

Matched rules	Final type (consensus)	n
Academic Researcher; Public Communicator	Academic Researcher	1
Industry Mover; Academic Researcher	Industry Mover	3
Industry Mover; Policy Shaper; Public Communicator	Policy Shaper	1
Industry Mover; Public Communicator	Public Communicator	1
Policy Shaper; Academic Researcher	Academic Researcher	1
Policy Shaper; Academic Researcher	Policy Shaper	1
Policy Shaper; Academic Researcher; Public Communicator	Academic Researcher	3
Policy Shaper; Public Communicator	Academic Researcher	1
Policy Shaper; Public Communicator	Policy Shaper	2
Policy Shaper; Public Communicator	Public Communicator	1
No rule matched	Public Communicator	1

Records were adjudicated independently by both coders and resolved in consensus according to the dominant logic of influence in the full record, with a documented rationale per case (audit trail).

As a descriptive check of the framework's homogeneity criterion, the mean pairwise simple-matching similarity across the twelve active indicators was computed within and between the final types: within-type similarity was 0.78 against a between-type similarity of 0.62, closely matching the corresponding values of the best purely statistical partition (0.81 and 0.61). Persona-level results were subsequently described using descriptive statistics only, in keeping with the complete-population design.



Supplementary Figs. S4

Supplementary Table S5. Prevalence of the twelve active indicators per persona (n = 57).

Indicator	Embedded Clinician	Academic Researcher	Public Communicator	Policy Shaper	Industry Mover
Med-tech industry role	0%	0%	9%	14%	100%
Teaching	0%	31%	27%	14%	0%
Further AI education	16%	31%	18%	0%	43%
Structured programme	0%	8%	27%	29%	14%
Individual research fellowship	11%	69%	36%	29%	57%
Publications: research	5%	85%	9%	57%	29%
Publications: opinion/debate	5%	46%	64%	14%	0%
Publications: knowledge	0%	31%	45%	14%	29%
Expert speaking	16%	23%	64%	57%	29%
Written opinion sharing	11%	0%	55%	43%	14%
Policy engagement	0%	38%	9%	100%	0%
Research activities	0%	46%	0%	0%	29%

Percentages are the share of persona members with at least one item in the category. Passive indicators (post-fellowship clinical activity, influencer activity, med-tech founding) and cohort membership are not shown; see the main text for cohort distributions.

S4 GRAMMS Reporting

The Design of the study is an explanatory sequential mixed methods (containing a quantitative digital-footprint analysis and a survey, followed by qualitative interviews).

#	GRAMMS criterion	How and where it is addressed in the manuscript	Reported
1	Justification describe the justification for using a mixed methods approach to the research question.	Introduction: frames impact measurement of AI fellowships as unresolved <i>“the measurement of impact is recognised as methodologically challenging yet essential”</i> ; motivating a combined approach. Methods › Study design: a <i>“retrospective mixed-methods evaluation ... [with] two complementary components”</i> , where a quantitative typology answers RQ1 (which skills alumni draw upon) and qualitative interviews answer RQ2 (how participation shapes development and influence). The Abstract restates the design.	Yes
2	Design describe the design in terms of the purpose, priority and sequence of methods.	Methods › Semi-structured interviews: <i>“Using a sequential mixed methods design, semi-structured interviews were conducted to provide additional context and explanation for the quantitative findings”</i> an explanatory sequential design. Purpose: explanation and complementarity (qualitative explains quantitative). Priority: quantitative-led; the footprint typology is primary and the survey has a <i>“secondary, clearly delimited role”</i> . Sequence: quantitative (footprint + survey) → qualitative (interviews).	Yes
3	Methods describe each method in terms of sampling, data collection and analysis.	Quantitative (Methods › Settings and Participants; Type Development). Sampling: complete population of cohorts 1 to 3 (n = 57), no sampling applied. Collection: systematic digital footprint (Google Scholar, PubMed, ResearchGate, LinkedIn, programme website, Google) plus a cross-sectional survey (Microsoft Forms, 20/57). Analysis: Kluge four-stage type construction; multiple correspondence analysis + hierarchical clustering with bootstrap / leave-one-out / permutation checks; descriptive statistics in Python 3.10. Qualitative (Methods › Semi-structured interviews). Sampling: purposive across the five personas and three cohorts. Collection: virtual interviews (ca. 20 min) via MS Teams, recorded, transcribed, de-identified. Analysis: inductive thematic analysis (Braun & Clarke) in NVivo 14, double-coded with consensus.	Yes
4	Integration describe where the integration has occurred, how it has occurred and who has participated in it.	Where (data level): Methods › Persona Development the survey verifies and completes the footprint records before indicator coding (stage 1) and survey free-text informs qualitative type construction (stage 3); cross-validation is quantified in Results <i>“Survey responses (20/57) corroborated footprint-</i>	Yes

		derived career information in all 20 cases”. Where (interpretation level): the explanatory sequential design (Methods › Semi-structured interviews) samples interviews from, and to explain, the footprint/survey patterns; the Discussion then brings all three strands together, using the qualitative themes to account for skills being mobilised alongside rather than within routine clinical care. How: cross-validation and triangulation combined with explanatory sequencing and complementarity. Who: KW and SP jointly performed extraction, independent coding and consensus adjudication across strands.	
5	Limitations describe any limitation of one method associated with the presence of the other method.	Discussion › Limitations. The footprint captures only publicly visible activity <i>“absence from the digital footprint is not evidence of absence”</i> is a blind spot (local, invisible contributions) that the survey and interviews are used to offset. Conversely those methods are self-reports: the survey has 35% response and the interviews are <i>“subject to self-selection and to attribution bias”</i>, so each method’s weakness is explicitly weighed against the others.	Yes
6	Insights describe any insights gained from mixing or integrating the methods.	Results and Discussion. Integration produced findings unreachable by either method alone: the survey shows skills are applied <i>“outside direct clinical practice, for reasons the survey itself cannot resolve”</i>, which the interviews then explain through organisational readiness and career structure. Combining strands also revealed that Embedded Clinicians, invisible in the footprint, exert real local influence <i>“The value of an embedded expert lies partly in being available, which no visibility-based metric will capture.”</i>	Yes

S5 COREQ

COREQ Reporting for the Publication *Mobilising Clinical AI Skills: A Mixed-Methods Evaluation of a National Medical Professional Upskilling Programme*.

#	COREQ item	How and where it is addressed in the manuscript	Reported on page no.
Domain 1: Research team and reflexivity			
1	Interviewer or facilitator which author conducted the interview?	Interviews were “conducted by either KW or SP” (Methods › Semi-structured interviews).	
2	Credentials interviewer’s credentials.	Stated only as “both ... trained in qualitative research methods”; Both Interviewer are PhD students, having done semi-structured interviews in their PhDs.	
3	Occupation at time of study.	KW and SP are both PhD students.	
4	Gender of the interviewer.	Both interviewer are female.	
5	Experience and training in qualitative methods.	“both of whom are trained in qualitative research methods” (Methods › Semi-structured interviews).	
6	Relationship established prior to study start?	None. Participants are recruited via the alumni network.	
7	Participant knowledge of interviewer what they knew about the researcher.	We stated the purpose of the interview in the invitation. Other than that no prior knowledge.	
8	Interviewer characteristics bias, assumptions, interests (reflexivity).	We explained the multidisciplinary background of the research team, including the mixture of clinical and non-clinical researchers.	
Domain 2: Study design			
9	Methodological orientation and theory.	“Thematic analysis ... conducted inductively following Braun and Clarke” (Methods › Semi-structured interviews).	
10	Sampling how were participants selected?	“selected purposively, aiming for balanced representation across the five personas and the three cohorts”.	
11	Method of approach how were they approached?	Email to consenting survey respondents plus a message in the alumni WhatsApp group (Methods › Semi-structured interviews).	
12	Sample size how many participated?	“A total of seven semi-structured interviews”	
13	Non-participation how many refused or dropped out, and why?	Not reported.	
14	Setting of data collection.	“Interviews were held virtually” via Microsoft Teams.	
15	Presence of non-participants anyone else present?	Not reported (one-to-one virtual interviews).	
16	Description of sample key characteristics.	“seven alumni from across all three cohorts with each persona represented at least once.”	

17	Interview guide questions, prompts; pilot tested?	Guide development described (by KW, from RQs and survey findings). Questions and prompts not provided to participants and pilot testing not stated.	
18	Repeat interviews carried out?	None	
19	Audio or visual recording.	<i>"All interviews were audio-recorded using Microsoft Teams".</i>	
20	Field notes made during and/or after?	None.	
21	Duration of interviews.	<i>"lasted approximately 20 minutes".</i>	
22	Data saturation discussed?	<i>"Saturation was reached by the 5th interview, as subsequent transcripts produced no new themes."</i>	
23	Transcripts returned to participants for comment or correction?	Not complete for pragmatic and timing related reasons.	
Domain 3: Analysis and findings			
24	Number of data coders	Two: each transcript <i>"independently coded by one researcher, with a second researcher reviewing the coding"</i> (KW, SP).	
25	Description of coding tree	A <i>"developing codebook"</i> is described but no coding tree is presented; consider a brief structure or Supplementary.	
26	Derivation of themes in advance or from data?	Inductive: codes were <i>"organis[ed] ... into themes that reflected participants' accounts"</i> .	
27	Software used to manage data.	<i>"NVivo Version 14".</i>	
28	Participant checking feedback on findings?	None.	
29	Quotations presented to illustrate themes; identified?	<i>"Illustrative quotations ... attributed using participant identification numbers";</i> Results quote IDs throughout (e.g., ID2–ID13).	
30	Data and findings consistent.	Each theme in Results is evidenced by supporting quotations consistent with the interpretation.	
31	Clarity of major themes.	Four major themes are clearly presented as Results subheadings.	
32	Clarity of minor themes diverse or minor cases.	Some within-theme variation is described, but minor themes or divergent cases are not separately delineated.	

S6 Survey Questions

Alumni Survey

The survey was designed to address RQ1 ("How do AI skills developed through the NHS Fellowship in Clinical AI influence alumni's professional development and contribute to their work within and beyond the NHS?") by supplementing findings from the digital footprint analysis. The survey covered demographics, continued use of fellowship-acquired AI skills, settings and levels of AI involvement, knowledge-transfer, research, education and voluntary roles, and perceived career impact. The responses were a mix of multiple-choice items, Likert-type scales and free-text responses.

Survey questions:

Questions 1-3 provided project and consent information, all survey participants were required to read and complete the consent questions to proceed with the survey questions.

4) Which cohort of the fellowship were you part of?

Option 1: Cohort 1 (2022-2023)

Option 2: Cohort 2 (2023-2024)

Option 3: Cohort 3 (2024-2025)

5) Have you continued to work in the NHS since completing the Clinical AI Fellowship?

Option 1: Yes- full time in the NHS

Option 2: Yes- part time in the NHS

Option 3: No

6) If you also work outside the NHS, what setting is this? (select all that apply)

Option 1: Research

Option 2: Medical Technology (MedTech) or digital health company

Option 3: Consultancy / advisory

Option 4: Other private sector (non-health)

Option 5: I do not work outside the NHS

Option 6: Other: [free text to provide detail]

7) What is your current role?

[Free text response]

8) Do you use the skills gained from the Clinical AI Fellowship in your current work?

Option 1: Yes

Option 2: No

9) In which setting do you **mainly** use these skills? (select up to max of 3)

Option 1: Clinical practice

Option 2: Academic research

Option 3: Education, teaching or training

Option 4: Industry (e.g. clinical advisor/CMO, clinical scientist, consultancy)

Option 5: Policy, governance, and leadership (public sector)

Option 6: Other: [free text to provide detail]

10) What best describes your **main** level(s) of involvement in AI? (select up to max of 3)

Option 1: User: using AI tools or applying general understanding of AI systems and limitations

Option 2: Embedder: supporting implementation, evaluation, or safe adoption of AI tools

Option 3: Creator: developing, training, or technically modifying AI models

Option 4: Driver: advising on or leading AI adoption, governance, or change

Option 5: Shaper: influencing AI strategy, policy, or system-level direction

Option 6: None. I am in no way involved with AI

11) Have you shared Clinical AI knowledge or skills with others since the Fellowship?

Option 1: Yes

Option 2: No

12) Where has this knowledge transfer occurred? (select all that apply)

Option 1: Within my day-to-day clinical role

Option 2: Within the NHS (e.g. panels, committees, teaching)

Option 3: Outside the NHS (e.g. panels for royal colleges, public media)

Option 4: Academic roles / research

Option 5: Industry work / collaborations

Option 6: Other: [free text to provide detail]

13) Are you involved in AI-related research projects?

Option 1: Yes

Option 2: No

14) In what **main** role are you involved in AI-related research?

Option 1: Project lead

Option 2: Co-investigator

Option 3: Contributor

Option 4: Other: [free text to provide detail]

15) Do you currently engage in any voluntary roles related to Clinical AI?

Option 1: Yes

Option 2: No

16) Please provide a short description of your volunteering work:

[Free text response]

17) Have you participated in another digital health or innovation programme?
(e.g. educational programmes in a cohort of students)

- Option 1: NHS Topol Digital Fellowship
- Option 2: NHS Clinical Entrepreneur Programme
- Option 3: BiteLabs Healthtech and Innovation Fellowship
- Option 4: Further tertiary education (e.g. PGDip, MSc, PhD)
- Option 5: No
- Option 6: Other: [free text to provide detail]

18) Have you been involved in any individual research fellowships related to AI / digital health? (e.g. Clinical AI Research Fellow at KCL)

- Option 1: Yes
- Option 2: No

19) What fellowship was this? (name and institution)

[Free text response]

20) Overall, how has the Clinical AI Fellowship influenced your career?

- Option 1: No impact
- Option 2: Some impact
- Option 3: Significant impact
- Option 4: Career-defining impact

21) Please provide additional comments to describe the impact of the Clinical AI Fellowship on your career

[Free text response]

22) Do you agree to be contacted again in relation to follow up questions or activities regarding this project?

- Option 1: Yes
- Option 2: No

Interview Guide

The semi-structured interview was designed to address RQ2 ("How do AI skills developed through the NHS Fellowship in Clinical AI influence alumni's professional development and contribute to their work within and beyond the NHS?").

The interview guide explored participants' motivations for undertaking the Fellowship, its perceived impact on their professional development and career opportunities, experiences of applying AI knowledge and skills in practice, barriers and enablers to skills utilisation, future career aspirations, and perceived support needs. Example questions are provided below. Additional prompts and follow-up questions were used as appropriate to explore topics in greater depth.

Example questions:

1. MOTIVATION: *What motivated you to join the Fellowship program?*
2. PROFESSIONAL IMPACT: *Do you feel the Fellowship has impacted your professional development or opportunities?*
 - *Has there been a change in your responsibilities?*
 - *Has the Fellowship opened any doors for you that would otherwise be inaccessible?*
3. FUTURE PLANS: *Thinking about your future career steps, where do you see yourself working in the next years?*
 - *Is there a clear career progression path for you within or beyond your organisation?*
4. ENABLERS: *What has helped you develop your skills and succeed in roles related to AI or digital transformation?*
 - *What aspects of the Fellowship were most impactful?*
5. CHALLENGES: *What are the main challenges you have encountered so far when trying to apply or further develop your AI skills?*
 - *Can you think of a time you wanted to use your AI-skills but couldn't?*
6. CURRENT GAPS: *What do you think would have helped/help you further in your role and future career aspirations?*

7. OTHER COMMENTS: *Are there any other comments you think are relevant to this conversation?*